



Perceptrons

Lecture 15 of “Mathematics and AI”



Outline

1. A neural approach to AI

Theoretical neuron, synapses, action potential, threshold voltage, plasticity

2. Perceptron

3. Learning algorithms

Perceptron learning, stochastic gradient descent, mini-batch

4. Variants of stochastic gradient descent

Momentum, Averaging, RMSProp, Adam



A neural approach to AI

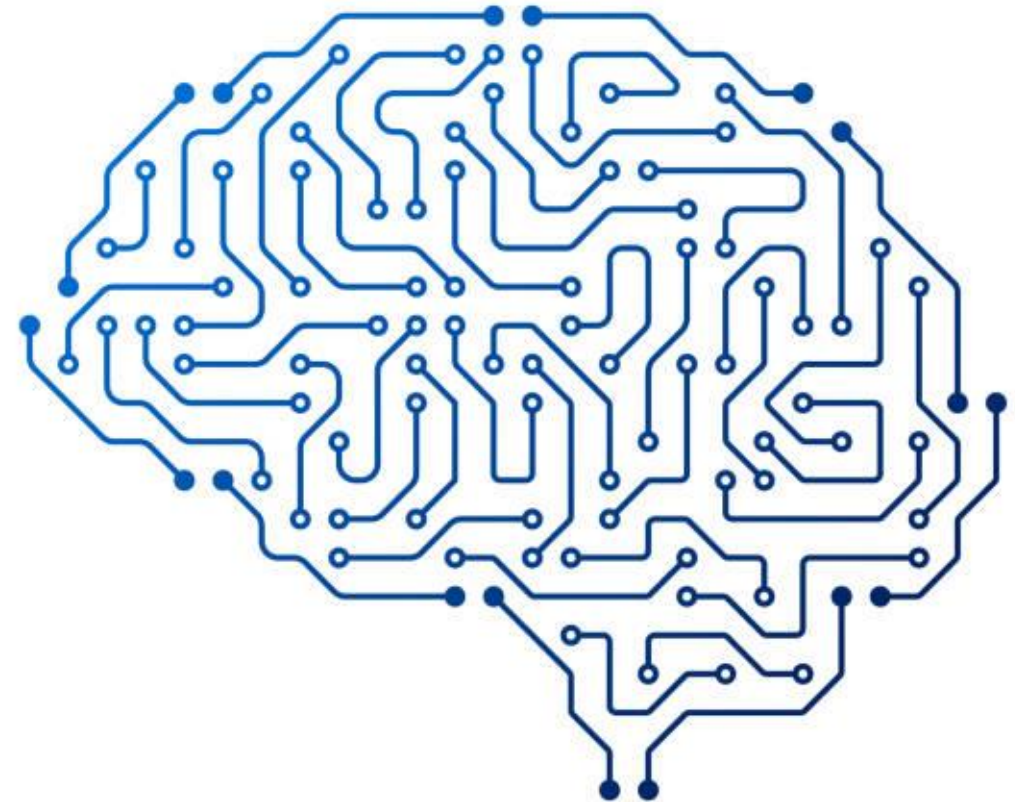
Back to beginnings ...

- AI: an interdisciplinary challenge
 - Mathematics,
 - physics,
 - philosophy,
 - psychology,
 - neuroscience



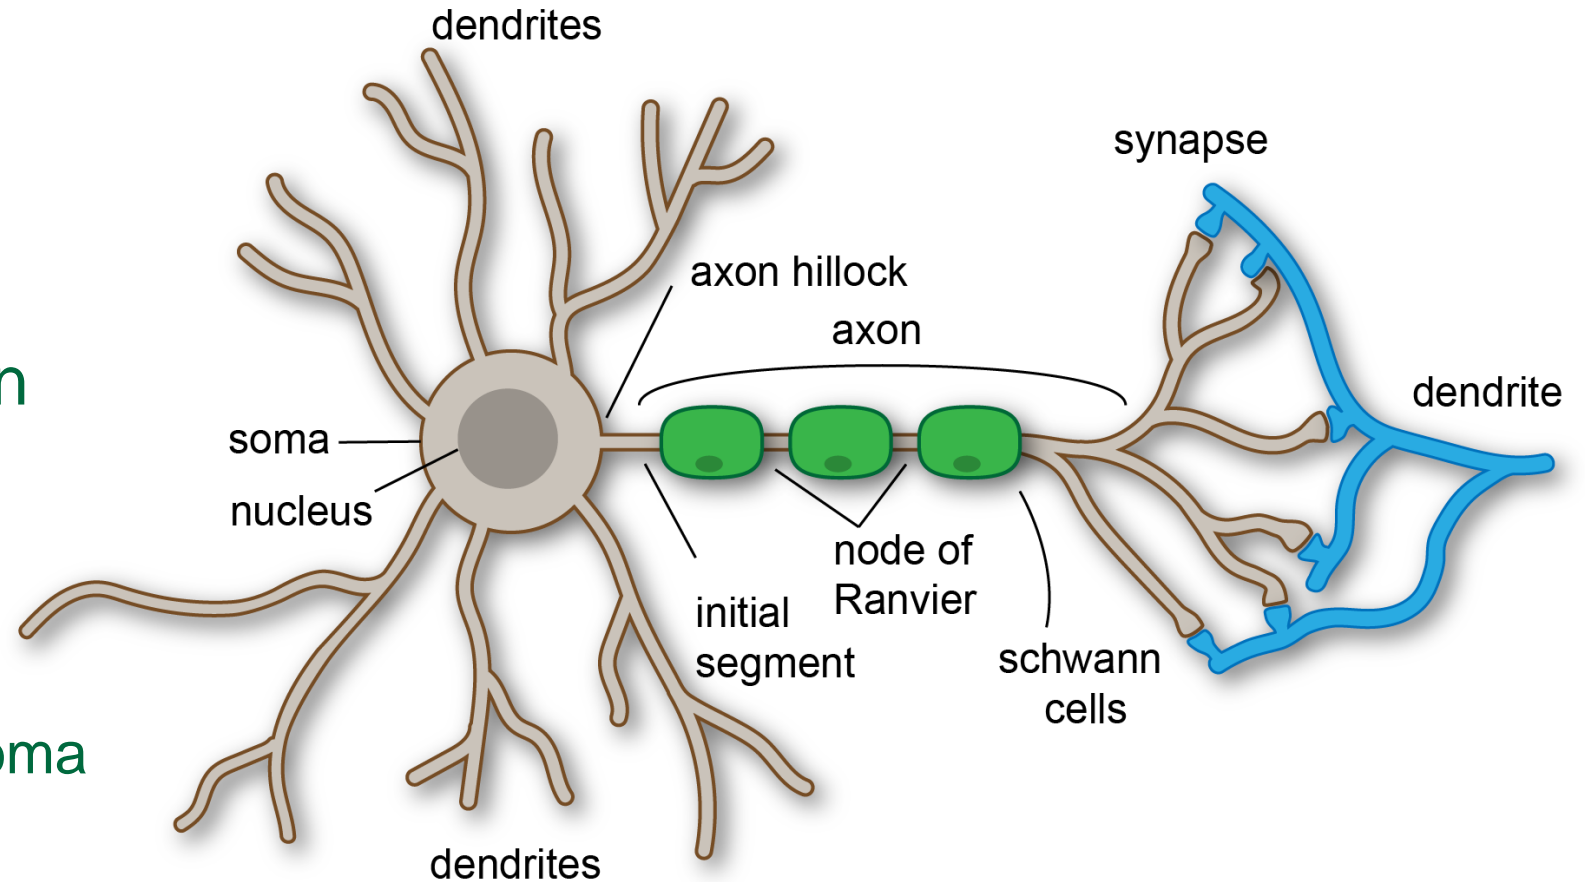
The human brain as a computer

- Is the brain as a *thinking machine*?
- How could one build such a machine from scratch?



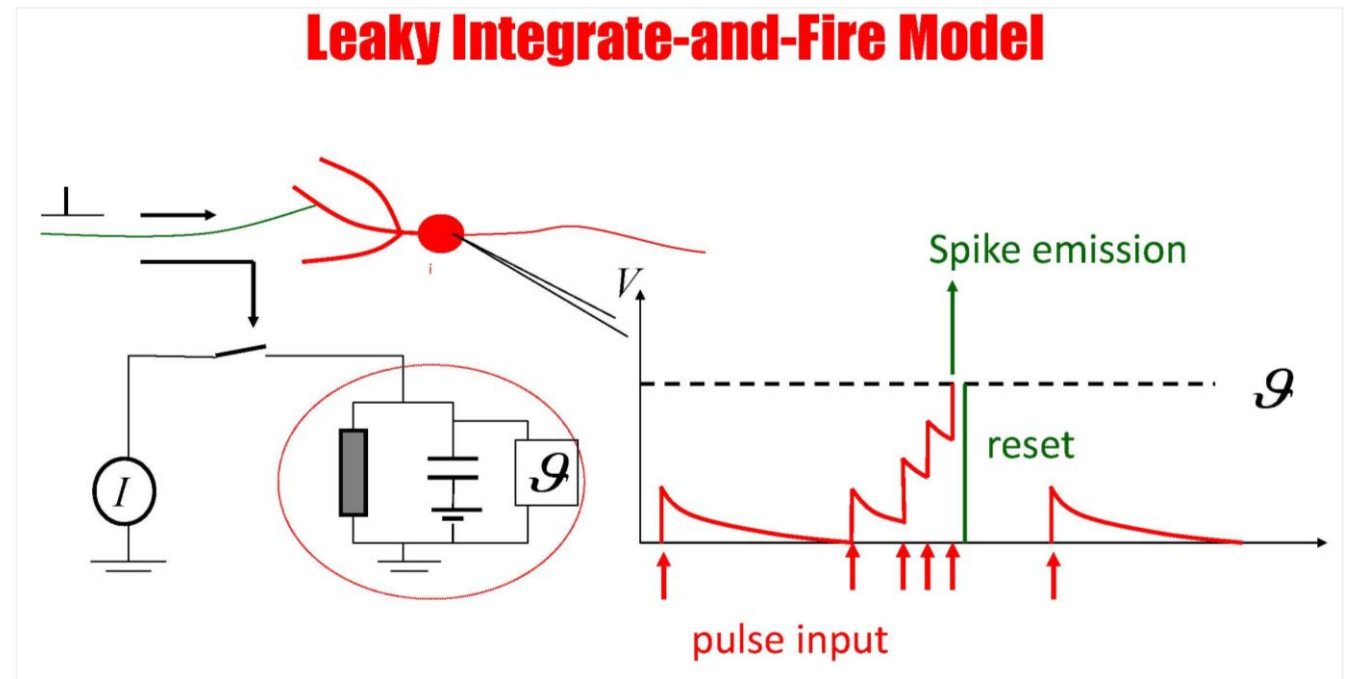
The neuron

- Fundamental information processing unit of the brain
- Neuron as I/O system:
 - Input: dendrites
 - Information aggregation: soma
 - Output: axon terminals
 - Communication: synapses



Theoretical models of neural dynamics

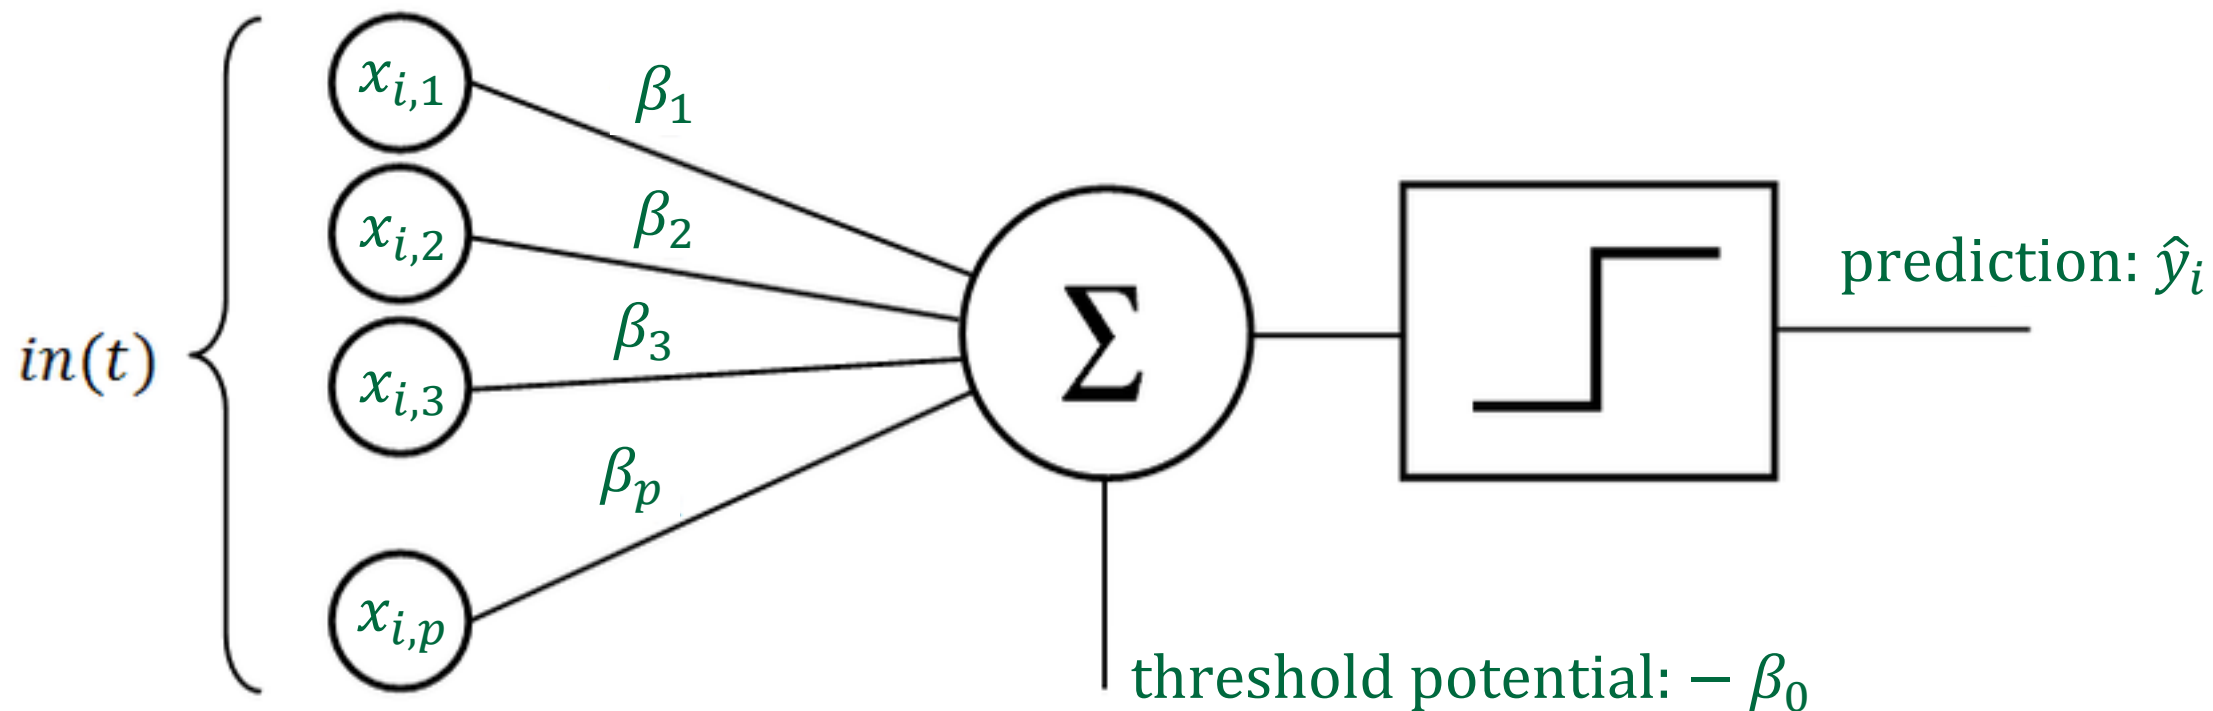
- Mathematical models of neurons
 - Hodgkin-Huxley model, Integrate-and-fire (IF) model, leaky integrate-and-fire (LIF) model
- Machine learning with spiking neural networks (SNNs)
 - Computer vision, robotics, ...



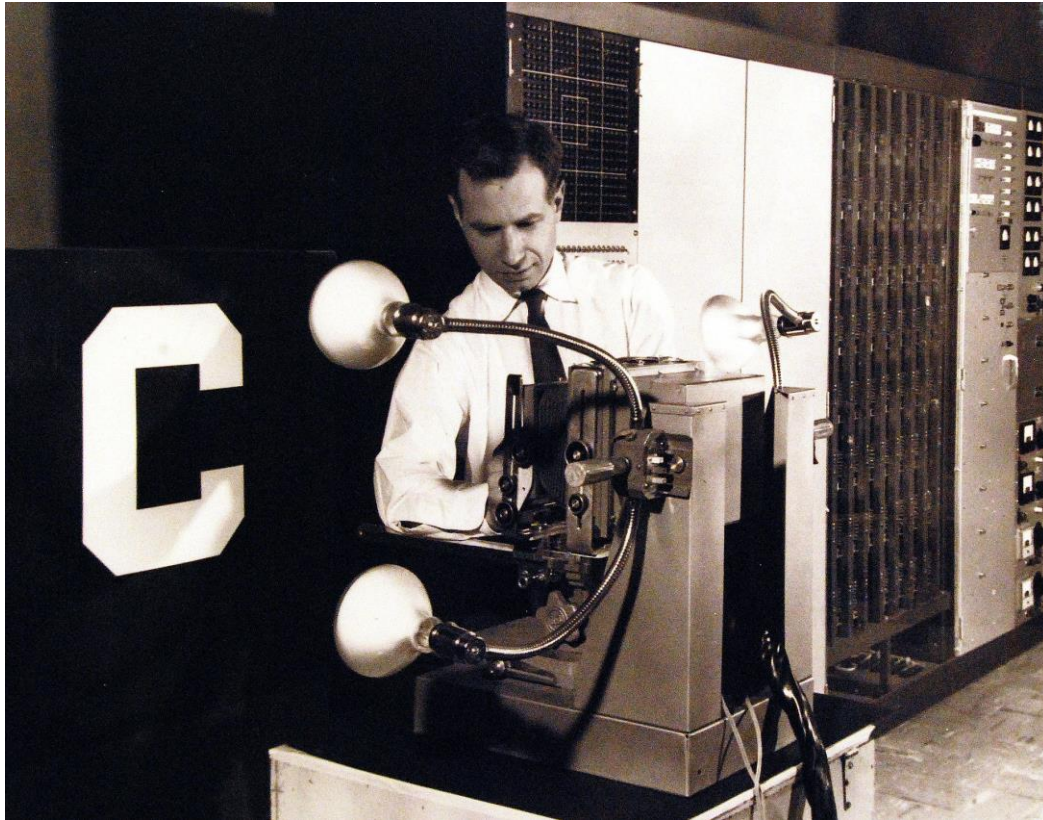


Perceptron

Perceptron



Perceptron



The Mark I Perceptron

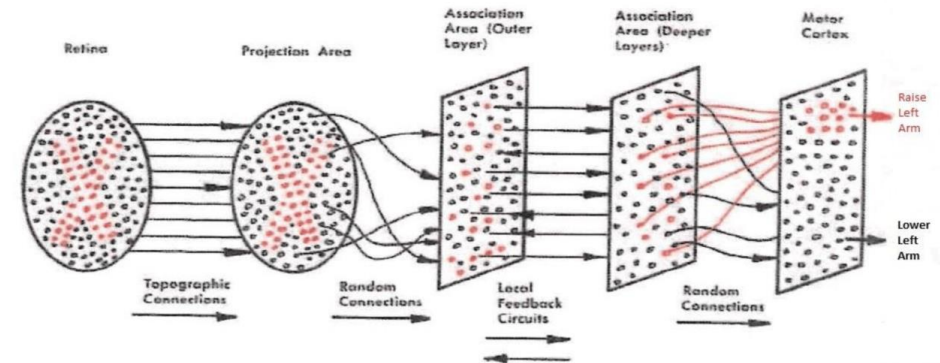


FIG. 1 — Organization of a biological brain. (Red areas indicate active cells, responding to the letter X.)

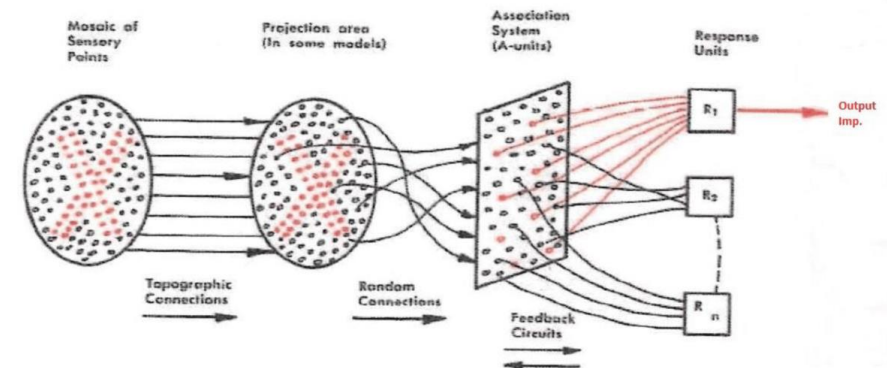
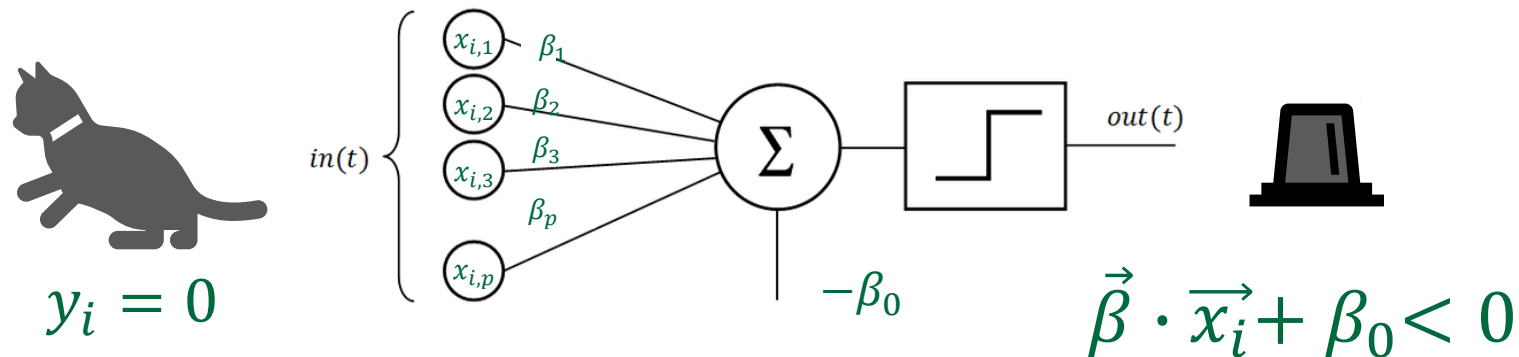
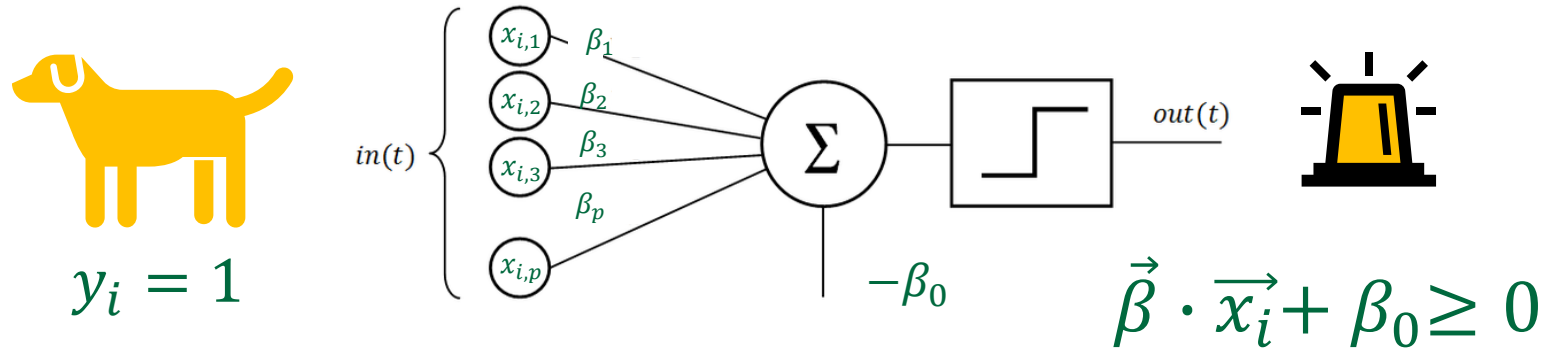


FIG. 2 — Organization of a perceptron.

Perceptron as a binary linear classifier



Linear decision boundary for binary classification

$$\{\vec{x}_i \mid \beta_0 + \vec{\beta} \cdot \vec{x}_i < 0\}$$



➤ \vec{x}_i on one side of hyperplane

➤ (\vec{x}_i, y_i) is “cat”

➤ $y_i = -1$

$$\{\vec{x}_i \mid \beta_0 + \vec{\beta} \cdot \vec{x}_i > 0\}$$



➤ \vec{x}_i on other side of hyperplane

➤ (\vec{x}_i, y_i) is “dog”

➤ $y_i = +1$

For **all** training samples:



$$y_i (\beta_0 + \vec{\beta} \cdot \vec{x}_i) > 0$$





Learning algorithms

Perceptron learning algorithm

- Natural learning is sequential
- Gradient descent does not work with step functions
- Perceptron learning algorithm

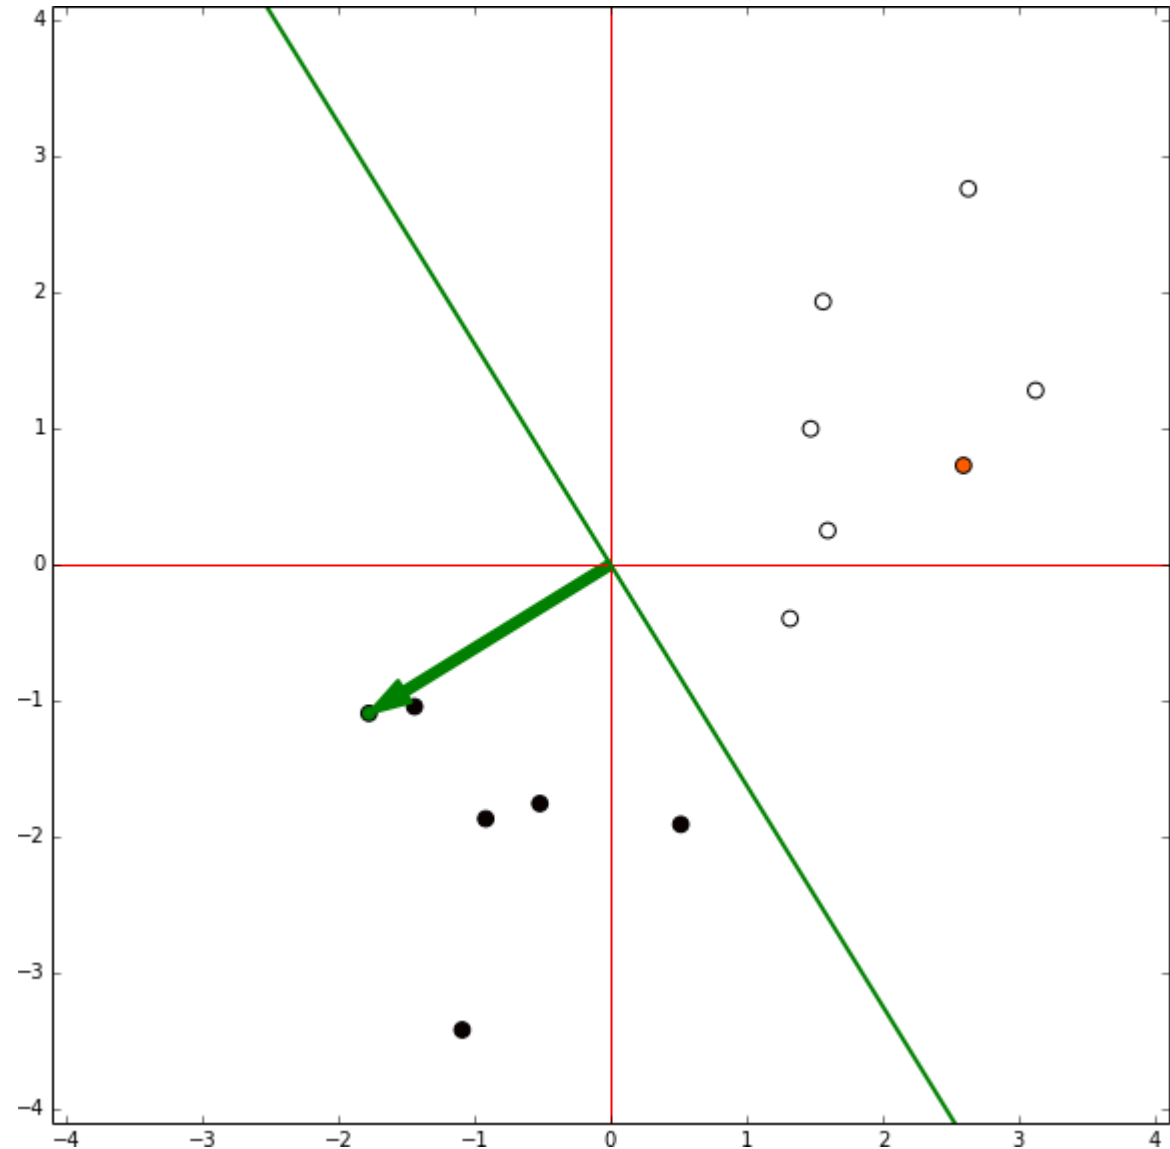
Algorithm: Perceptron Learning Algorithm

```
 $P \leftarrow \text{inputs with label } 1;$   
 $N \leftarrow \text{inputs with label } 0;$   
Initialize  $\mathbf{w}$  randomly;  
while !convergence do  
    Pick random  $\mathbf{x} \in P \cup N$  ;  
    if  $\mathbf{x} \in P$  and  $\mathbf{w} \cdot \mathbf{x} < 0$  then  
        |  $\mathbf{w} = \mathbf{w} + \mathbf{x}$  ;  
    end  
    if  $\mathbf{x} \in N$  and  $\mathbf{w} \cdot \mathbf{x} \geq 0$  then  
        |  $\mathbf{w} = \mathbf{w} - \mathbf{x}$  ;  
    end  
end
```

//the algorithm converges when all the inputs are classified correctly

Perceptron learning

- Natural learning is sequential
- Gradient descent does not work with step functions
- Perceptron learning algorithm
 - + Works for with step functions
 - Converges slowly or not at all



Stochastic gradient descent (SGD)

- SGD: Sequential gradient-based learning

Gradient descent (full batch)	Stochastic gradient descent
<p>Loss function:</p> $L_{GD}(\vec{\beta}) = \frac{1}{n} \sum_{i=1}^n \ell_i(\vec{\beta}) = \frac{1}{n} \sum_{i=1}^n \ell_i(\vec{\beta}, \vec{x}_i, y_i)$ <p>Update rule:</p> $\vec{\beta}' = \vec{\beta} - \gamma \nabla L_{GD}(\vec{\beta}) = \vec{\beta} - \frac{\gamma}{n} \sum_{i=1}^n \nabla \ell_i(\vec{\beta})$	<p>Loss function:</p> $L_{SGD}(\vec{\beta}) = \ell_i(\vec{\beta}, \vec{x}_i, y_i)$ <p>Update rule:</p> $\vec{\beta}' = \vec{\beta} - \nabla L_{SGD}(\vec{\beta}) = \vec{\beta} - \gamma \nabla \ell_i(\vec{\beta})$



Stochastic gradient descent (SGD)

- Example: OLS with one feature

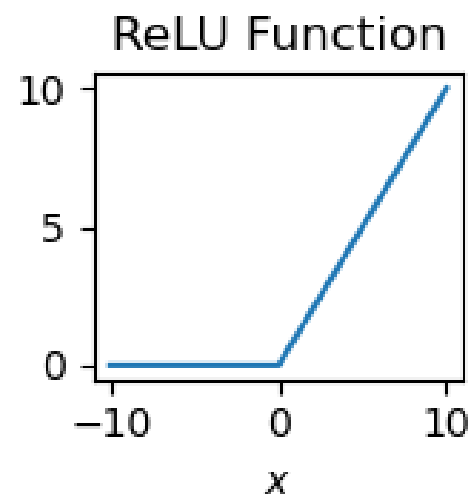
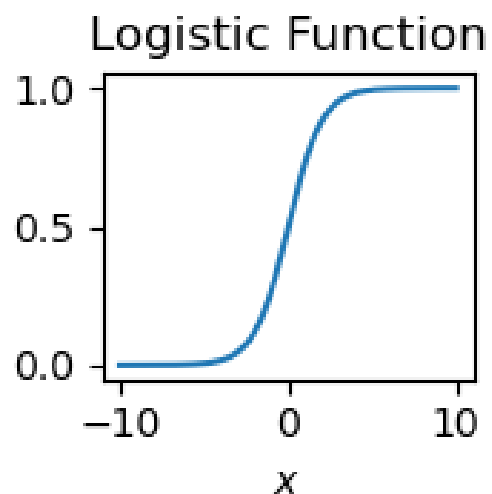
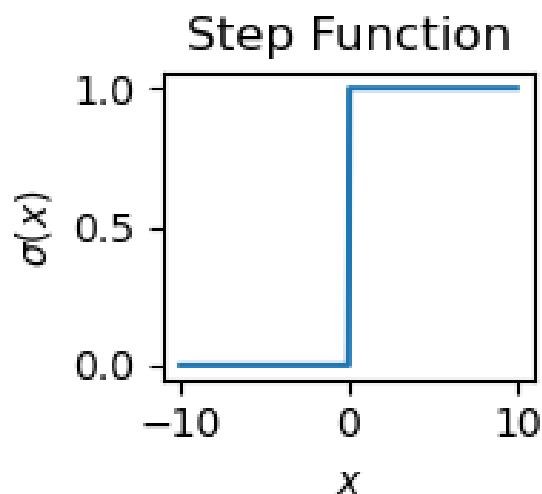
$$L_{GD}(\vec{\beta}) = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2 = \frac{1}{n} \sum_{i=1}^n (\beta_0 + \vec{\beta} \cdot \vec{x}_i - y_i)^2$$

$$L_{SGD}(\vec{\beta}) = \ell_i(\vec{\beta}, \vec{x}_i, y_i) = (\beta_0 + \vec{\beta} \cdot \vec{x}_i - y_i)^2$$

$$\vec{\beta}' = \vec{\beta} - 2\gamma(\beta_0 + \vec{\beta} \cdot \vec{x}_i - y_i)\vec{x}_i$$

Activation functions

- SGD does not work with step functions, but with many other naturalistically motivated activation functions
- Examples: logistic/sigmoid function, ReLU (Rectified linear unit)





Linear activation function?

- Unbounded function:
 - Not biologically motivated
 - Not well suited for classification problems
- Linear function:
 - $\vec{\beta} \cdot \vec{x}_i$ is already a linear transformation (LT)
 - $\text{LT}(\text{LT}(\vec{x}_i))$ is just another LT on \vec{x}_i
 - Does not change model expressiveness



Variants of SGD

Variants of SGD



GD can get stuck in $\nabla L_{GD}(\vec{\beta}) = 0$ zones



SGD converges slowly and is sensitive to outliers



➤ Variants of SGD that include smoothening effect

➤ Averaging / mini-batch

➤ Momentum

➤ RMSProp

➤ Adam

Averaging/ mini-batch learning



- Split training data into “mini-batches” of size m
- Loss:

$$L_{SGD}(\vec{\beta}) = \frac{1}{m} \sum_{i=1}^m \ell_i(\vec{\beta}, \vec{x}_i, y_i)$$

- Update rule:

$$\vec{\beta}' = \vec{\beta} - \gamma \nabla L_{SGD}(\vec{\beta})$$

- Improves robustness to outliers

SGD with momentum



➤ Add momentum to how $\vec{\beta}$ changes

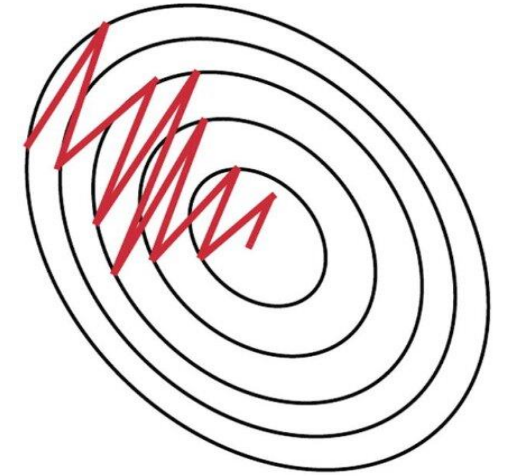
➤ Loss:

➤ Standard SGD loss or mini-batch

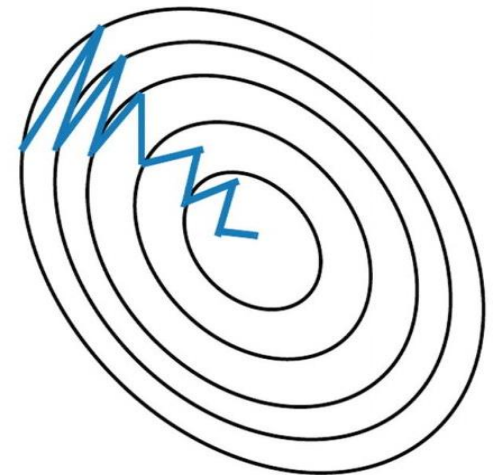
➤ Update rule:

$$\vec{\beta}' = \vec{\beta} + \Delta\vec{\beta}' \quad \text{with} \quad \Delta\vec{\beta}' = -\gamma \nabla L_{SGD}(\vec{\beta}) + \alpha \Delta\vec{\beta}$$

➤ Improves robustness for high-variance data



Stochastic Gradient
Descent **without**
Momentum



Stochastic Gradient
Descent **with**
Momentum

RMSProp, Adam



- RMSProp (Root mean square propagation)
- Adjust learning rate for each parameter

- Update rule:

$$\vec{\beta}' = \vec{\beta} + \Delta\vec{\beta}' \text{ with } \Delta\vec{\beta}' = -\frac{\gamma}{v} \nabla L_{SGD}(\vec{\beta}) \text{ and } v \text{ moving average of } \Delta\vec{\beta}'$$

- Improves robustness to high-variance features
- Adam (Adaptive Moment Estimation)
 - combines RMSProp with momentum